

# Thread-Safe Cache Study Guide

## Overview

The **Thread-Safe LRU Cache** demonstrates advanced concurrency patterns for high-performance caching systems with **reader-writer synchronization** and **lazy LRU management**.

## Key Concepts Covered

### 1. Reader-Writer Synchronization

```
std::shared_mutex mutex_;

// Multiple concurrent readers
std::shared_lock<std::shared_mutex> read_lock(mutex_);

// Exclusive writer access
std::unique_lock<std::shared_mutex> write_lock(mutex_);
```

### 2. LRU (Least Recently Used) Algorithm

```
// Doubly-linked list + hash map = O(1) operations
std::unordered_map<Key, Iterator> cache_map_;
std::list<std::pair<Key, Value>> lru_list_;
```

### 3. Lazy vs Eager Updates

- **Eager:** Update LRU order on every read (high contention)
- **Lazy:** Update LRU order only on writes (better performance)

## Data Structure Design

### Hybrid Architecture

```
struct CacheNode {
    Key key;
    Value value;
    // Implicit position in linked list
};

// Fast lookup: O(1)
std::unordered_map<Key, list_iterator> map_;

// LRU ordering: O(1) insert/remove
std::list<CacheNode> lru_list_;
```

## Operations Complexity

Operation	Time	Space	Contention
<b>get()</b>	O(1)	O(1)	Shared lock
<b>put()</b>	O(1)	O(1)	Exclusive lock
<b>eviction</b>	O(1)	O(1)	During put()

## Real-World Applications

### System Caches

- **CPU caches:** L1/L2/L3 cache hierarchies
- **Database buffer pools:** Page caching
- **Web caches:** HTTP response caching
- **CDNs:** Content delivery networks

### Application Caches

- **Redis/Memcached:** Distributed caching
- **In-memory caches:** Application-level caching
- **File system caches:** OS page cache
- **Browser caches:** Resource caching

## Interview Questions & Answers

**Q: “Why use shared\_mutex instead of regular mutex?”**

**A:** Performance optimization for read-heavy workloads: - **Concurrent reads:** Multiple threads can read simultaneously - **Exclusive writes:** Only one writer, no readers during writes - **Typical cache ratio:** 90% reads, 10% writes - **Speedup:** 5-10x improvement in read-heavy scenarios

**Q: “How do you handle cache eviction?”**

**A:** LRU eviction strategy:

```
if (cache_map_.size() >= capacity_) {  
    // Remove least recently used (back of list)  
    auto lru_key = lru_list_.back().first;  
    cache_map_.erase(lru_key);  
    lru_list_.pop_back();  
}
```

**Q: “What about thread safety during eviction?”**

**A:** Write lock protects entire operation: - Check capacity while holding exclusive lock - Eviction and insertion are atomic - No partial states visible to readers

**Q: “Why not update LRU on every read?”**

**A: Lazy LRU** trade-off: - **Pros:** Better performance, less contention - **Cons:** Slightly less accurate LRU ordering - **Real-world:** Most caches use approximations anyway

## Design Patterns Demonstrated

### 1. Adaptive Locking Strategy

```
// Readers use shared locks
std::shared_lock<std::shared_mutex> lock(mutex_);

// Writers use exclusive locks
std::unique_lock<std::shared_mutex> lock(mutex_);
```

### 2. RAII Lock Management

- Automatic lock release
- Exception safety
- Clear lock scope boundaries

### 3. Template Specialization

```
template<typename Key, typename Value>
class ThreadSafeCache {
    // Generic for any key-value types
    // Requires Key to be hashable
};
```

## Performance Optimization Techniques

### 1. Lock Granularity

- **Coarse-grained:** One lock for entire cache
- **Fine-grained:** Bucket-level locking (more complex)
- **Trade-off:** Simplicity vs scalability

### 2. Memory Layout

```
// Cache-friendly access patterns
std::list<std::pair<Key, Value>> lru_list_; // Sequential access
std::unordered_map<Key, Iterator> map_; // Random access
```

### 3. Move Semantics

```
void put(Key key, Value value) {
    // Move to avoid copies
```

```
lru_list_.emplace_front(std::move(key), std::move(value));
}
```

## Cache Replacement Policies

### LRU Alternatives

- **LFU (Least Frequently Used)**: Count-based eviction
- **FIFO**: Simple queue-based eviction
- **Random**: Minimal overhead
- **ARC (Adaptive Replacement Cache)**: Balances recency and frequency

### When to Use Each

- **LRU**: General-purpose, good temporal locality
- **LFU**: Frequency matters more than recency
- **FIFO**: Simple, predictable behavior
- **Random**: Very low overhead, surprisingly effective

### Test Scenarios Covered

1. **BasicOperations**: put/get functionality
2. **EvictionBehavior**: LRU ordering verification
3. **ConcurrentAccess**: Multiple readers/writers
4. **CapacityLimits**: Proper eviction handling
5. **MemoryLeaks**: Resource cleanup verification
6. **StressTest**: High contention scenarios

## Common Interview Challenges

### 1. “Implement LRU cache in 15 minutes”

**Key points to hit:** - Hash map for  $O(1)$  lookup - Doubly-linked list for  $O(1)$  eviction - Move recently used to front - Evict from back when full

### 2. “How would you scale this to multiple machines?”

**Distributed cache considerations:** - **Consistent hashing**: Distribute keys across nodes - **Replication**: Handle node failures - **Cache coherence**: Keep replicas synchronized - **Hot spotting**: Load balancing strategies

### 3. “What if the cache is too big for memory?”

**Hybrid storage strategies:** - **Tiered caching**: Memory + disk - **Compression**: Reduce memory footprint - **Eviction to disk**: Spillover storage - **Memory mapping**: OS-managed paging

## Production Considerations

### Monitoring & Metrics

- **Hit rate:** Percentage of cache hits
- **Eviction rate:** How often items are removed
- **Memory usage:** Cache size monitoring
- **Contention:** Lock wait times

### Configuration Tuning

- **Cache size:** Balance memory vs hit rate
- **Eviction policy:** Match access patterns
- **Concurrency level:** Readers vs writers ratio
- **Warmup strategy:** Pre-populate important data

## Advanced Extensions

### Lock-Free Cache

- **Hazard pointers:** Safe memory reclamation
- **Atomic operations:** CAS-based updates
- **Memory ordering:** Sequential consistency
- **Complexity:** Much harder to implement correctly

### Partitioned Cache

- **Segment locks:** Reduce contention
- **Hash-based partitioning:** Distribute load
- **NUMA awareness:** Thread-local caches

This cache implementation demonstrates sophisticated concurrent data structure design - a key skill for high-performance systems and distributed architectures!